Dataset of real words, but they come with real human ratings

We can run the script on the Engelthaler & Hills 4,997-word norms (github.com/tomasengelthaler/HumorNorms). Those are real words, but they come with real human ratings. It's a good second benchmark that shows your pipeline correlates with people.

https://github.com/tomasengelthaler/HumorNorms